E3: the weight distribution as a knob, direct loss -- paired gain over the lasso drawn at the same level
[laptop rehearsal: p = 10, true C = 2I, 4 replicates, 16 graphs per cell]

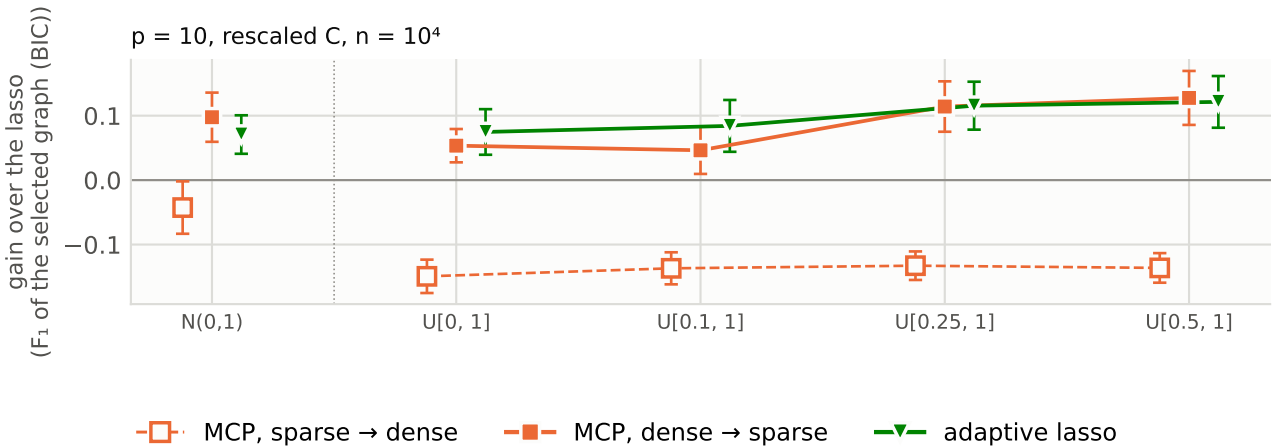

Graph per data set: the support that minimises loss + BIC penalty over the path (no truth used). Gain = mean over the cell's graphs of F_1(estimator) − F_1(lasso) on the same graph (same seed, same generator); bars = one standard error of that mean.